

# A Study of Cuckoo Hashing

Zhijia Chen

No Institute Given

**Abstract.** The abstract should briefly summarize the contents of the paper in 150–250 words.

**Keywords:** First keyword · Second keyword · Another keyword.

## 1 Introduction

It is said that the data are to this century what oil was to the last one: a driving force of growth and change. Indeed, data has been acting as fuels for many techniques such as machine learning, the Internet of Things and big data analysis, which in turn promote people to generate even more data. The data in last few decades has been growing exponentially and is supposed to keep the exploding trend. According to the report of International Data Corporation (IDC) published in 2014, the total size of data in the digital universe will reach to 44 zettabytes by 2020, and nearly 50% of cloud-based services will rely on data in storage systems. What's more, many computing devices, from small IoT sensors to powerful server in large scale data centers, are collecting and analyzing ever-greater amounts of data. A users in routine generate queries on hundreds of Gigabytes of data stored in their local storage or cloud storage systems, and commercial companies generally handle Terabytes and even Petabytes of data each day[SmartCuckoo].

Thus it is more and more challenging for cloud storage systems to serve queries in a real-time manner and keep a stable throughput over time, which consumes substantial resources to support query-related operations[]. And effective hashing scheme is becoming a critical factor for query service to keep up with the data growth. A hash table maps a keys to values and it uses a hash function to compute the position in the table where the corresponding data could be stored or found. Hash tables are so adorable for data indexing because they have a constant lookup complexity on average which supports fast query response whatever the data size is. But hash table has also has its drawback — hashing collision. Because a hash function usually maps multiple distinctive elements into one element, it may points multiple different keys to the same slot.

Thus one of the most basic tasks when designing a hashing scheme is to handle hashing collision.

Generally there are three strategies for resolving hashing collisions[]:

- **Perfect hashing:** Choose a hash function with no collisions. The idea of using hashing functions that prevent the collisions seems appealing. However,

it requires complete knowledge of all the keys to construct a perfect hashing function [fundamental study perfect hashing], which is not feasible for large cloud data storage centers that need to handle new coming data all the time.

- **Open addressing:** Allow items to overflow out of their target bucket and into other places. With open addressing, data collisions are resolved by finding alternative locations in hash table by probing a sequence of possible indices. Suppose  $T_i$  is the  $i$ th probing position and  $h(x), h_1(x), h_2(x)$  are hash functions, three basic probing methods are listed below.
  - Linear probing:  $T_i(x) = h(x) + i$
  - Quadratic probing:  $T_i(x) = h(x) + i^2$
  - Double hashing:  $T_i(x) = h_1(x) + i \times h_2(x)$
- **Closed addressing (separate chaining):** All the colliding items should be stored in the target bucket with the help of some auxiliary data structures. The closed addressing (also called separate chaining) resolves hashing collision by storing all colliding items in an auxiliary linked list or BST. The basic implementation are simple, but its performance is sensitive to data distribution, as a lookup in a bucket with collisions will involve a searching the auxiliary data structure, which could be very expensive in some unfortunate cases where too much keys are collided in the bucket due to skewed distribution of data.

Both open addressing and closed addressing, however, will suffer from great latencies when there are too many collisions in the target bucket under some unfortunate cases. So people are looking for hashing schemes that could have a constant lookup time even in the worst case, and cuckoo hashing is one of those candidates. Cuckoo hashing is a fast and simple hashing scheme with constant-time worst-case lookup ( $O(\ln \frac{1}{\epsilon})$ ) and has  $(1 + \epsilon)n$  memory consumption, where  $\epsilon$  is a small constant. It is widely used in large storage systems for its desirable properties. Unlike simple hashing schemes that provide only one bucket, cuckoo hashing provides  $N$  ( $N = 2$  in common practice) possible locations for each item by using  $N$  hash functions to reduce the change of collision. And to resolve collisions, old items are kicked out by the new comers and turn to their alternate buckets. To look up for an item, the cuckoo hashing only needs to search for a small number of buckets (depends on the number of hash functions being used), which is a constant time even under the worst case.

However, cuckoo hashing suffers from substantial performance penalty during insertion when the recursively kicking out operation turn out be an endless loop. The endless loop indicates a failure for insertion and need to perform a rehashing in order to accommodate all items. The endless loop itself, while useless for the result, consumes lots of time and resources, and it's only after exhaustive attempts the cuckoo hashing could realize that it's endless and perform rehashing, which leads to non-deterministic performance and is not desirable for cloud storage systems that are supposed to handle queries in real-time manner.

This paper presents the basic cuckoo hashing and some of its state-of-the-art variants that addressed on its non-deterministic performance issue during insertion. The rest of the paper is organized as follows. In section 2, we present

the basic cuckoo hashing and illustrate its operations with some examples. We describe three cuckoo hashing variants—cuckoo hashing with a stash (CHS), SmartCuckoo and horton tables—in the following three sections, and then we compare the three variants and conclude in section 6.

## 2 Standard Cuckoo Hashing


The standard cuckoo hashing scheme was proposed by Pagh and Rodler in [cuckoo hashing]. Suppose we attempt to accommodate  $n$  items with two tables,  $T_1$  and  $T_2$ , each has  $\frac{n}{2}$  buckets and one hash function ( $h_1$  for  $T_1$  and  $h_2$  for  $T_2$ ). Each bucket can store at most one item. To insert an item  $x$ , we place it in  $T_1[h_1(x)]$  if that bucket is empty. Otherwise (suppose it is taken by item  $a$ ), we kick out  $a$  in  $T_1[h_1(x)]$  and replace it with  $x$ , then we try to insert  $a$  into  $T_2[h_2(a)]$ . If  $T_2[h_2(a)]$  is empty, then we are done, and if it is not (suppose it is taken by item  $b$ ), then we further replace  $b$  with  $a$  and turn  $b$  to bucket  $T_1[h_1(b)]$  and so on, bouncing between  $T_1$  and  $T_2$  until all items stabilize.

Inserting item  $x$  will end up in one of the following three cases, as illustrated in Fig. 1. The directed edges between the two tables gives the potential positions of the items. Each edge points from one item's actually storage location to its alternative location.

- Item  $x$  is inserted without any kick-out operation. As shown in Fig. 1(a), two items ( $a$  and  $b$ ) are initially stored in  $T_1$  and  $T_2$ . When inserting  $x$ , its two candidate buckets are  $T_1[0]$  and  $T_2[5]$ .  $T_1[0]$  is occupied by item  $a$  while  $T_2[5]$  is empty, so we place  $x$  in  $T_2[5]$  and add an edge directed from  $T_2[5]$  to  $T_1[0]$ .
- Item  $x$  is inserted with finite kick-out operations. As shown in Fig. 1(b), items  $c$  and  $d$  are inserted into the hash tables before  $x$ . The two candidate bucket of  $x$  are occupied by  $a$  and  $d$  respectively. To accommodate  $x$ , we kick out  $a$  and move it to  $T_2[1]$ , which further moves  $b$  to  $T_1[2]$ . This procedure is repeated until an empty bucket is found at  $T_2[3]$ . The kick-out path is  $x \rightarrow a \rightarrow b \rightarrow c$ .
- Inserting item  $x$  runs into infinite kick-out operations. As shown in Fig. 1(b), items  $e$  is inserted into the hash tables before  $x$ . There is no vacant bucket to accommodate  $x$  even after substantial kick-out operations, which results in an endless loop. The cuckoo hashing has to carry out a rehashing operation[43].

A lookup or delete operation only needs to probe the two candidate buckets of an item, which requires a constant time. For example, to query  $x$ , we only need to check  $T_1[0]$  and  $T_2[5]$  as shown in Fig. 1.

If the two hash functions are chosen independently from an appropriate universal hash family, then we can expect to insert all  $n$  items successfully with at most  $\alpha \log n$  kick-out operations for any particular item with probability  $1 - O(1/n)$  [cuckoo hashing], and  $\alpha$  is a sufficiently large constant. Further more, if we limit the number of kick-out operation to  $\alpha \log n$  before rehashing, then



**Fig. 1.** The standard cuckoo hashing data structure.


the expected time to insert all  $n$  items successfully into the table is  $O(n)$ , which means the expected insert complexity of the standard cuckoo hashing is also a constant.[cuckoo]

### 3 Cuckoo Hashing with a Stash

Cuckoo hashing supports simple worst-case  $O(1)$  lookups, however, collisions are still troublesome for insertion. Though rare, there is always a chance that on the insertion of a new item, we could not find a place to hold the item even after substantial attempts, causing a failure. In such case, a full rehash is required to find new hash functions to accommodate all items appropriately. Theoretically speaking, the insertion failure of cuckoo hashing happens with a low probability ( $O(n^{-c})$  for some constant  $c \geq 1$ ). The rehashing has very limited impact on the average query performance, but the very existence of those unfortunate chances make it hard to guarantee the performance. Moreover, the constant  $c$  in  $O(n^{-c})$  failure probability bound is hard to be measured thus difficult to guarantee that  $c$  is above the desirable level in practice. For example, values of  $c \leq 3$  may lead too much failure of commercial applications.

Cuckoo hashing with a stash (CHS) is a variant of cuckoo hashing that utilizes a small amount of memory to construct a small stash outside the main hash tables. Specifically, by storing a constant number of items outside the table in an area that called stash, we can dramatically reduce the probability of insertion failure where a rehash operation is necessary. The intuition behind the approach is that, if items cause insertion failure independently, then we can expect the number of items that cause insertion failure to be  $P_r(S \geq s) = O(n^{-cs})$  for some constant  $c > 0$  and every constant  $s \geq 1$ . And we can identify those problematic items during insertion and store them in the stash, dramatically reducing the failure probability bound[CHS].

To detect the problematic item, CHS relies on *cuckoo graph* which is a multi-graph that derived from the two hash tables of cuckoo hashing. The graph corresponding to the Fig. 1 is shown in Fig. 2. In a cuckoo graph, each vertex



**Fig. 2.** The corresponding cuckoo graph of Fig. 1

represents an alternative bucket of some items (no matter the bucket is actually occupied or not), and each directed edge corresponds to an item and is directed from its actual stored bucket to its alternative bucket. Note that the cuckoo graph can only represent cuckoo hashing that has two hash functions. We can easily determine if an insertion would fail or not with the help of cuckoo graph. Since each vertex has at most one outgoing edge, each component in a cuckoo graph has at most one cycle. Inserting a new element will add a new edge in the graph, and if the target vertex is connected to a cycle, then the insertion is doomed to fail[Cuckoo Hashing: Further Analysis. Information Processing Letters]. During insertion, the CHS will put an item in the stash whose corresponding edge belongs to a cycle, thus avoids insertion failure. And the items in the stash will be inserted into the hash tables late whenever there are available buckets. So a rehash is only needed when the stash is full and none of the items in the stash could be successfully inserted into the tables without rehashing.

There are many ways to implement the insertion algorithm for CHS. One way is to track the bucket that has been visited during insertion. During a successful insertion, at most one vertex of the cuckoo graph is visited (kicking out the item in the corresponding bucket) more than once, and no vertex is visited more than twice. Thus we can remember the slots that has been visited during the insertion, and when a bucket has been visited more than twice, put the item into the stash. However, logging the buckets that has been visited could be expensive, thus the author suggests that we set a limit of  $\alpha \log n$  on the number of possible kick-out operations. And if  $\alpha \log n$  evictions do not suffice to settle down, then we 'roll back' to the original configuration and try to insert the new coming item second time with the 'cycle detection' mechanism[CHS].


[CHS] has proved that for every constant integer  $s \geq 1$ , for a sufficiently large constant  $\alpha$ , the size  $S$  of the stash after all items have been inserted satisfies  $P_r(S \geq s) = O(n^{-s})$ .

## 4 SmartCuckoo

Many cuckoo hashing variants have addressed the problem of endless kickout operations when countering an insertion failure. The common feature of them is to set an upbound on the number of evict operations before performing rehashing.

However, those kick-out loops executed in failure insertion consume substantial resources and they are just fruitless efforts. The insertion performance of cuckoo hashing will be much better if we can identify the endless loops and avoid wasting time and resources on them. With such motivation, *SmartCuckoo* is designed to predetermine and identify endless loops during insertion. Like *CHS*, *SmartCuckoo* also leverages on the cuckoo graph to identify the cycles. However, instead of remembering visited buckets to track the component of cuckoo graph that involved in current execution, *SmartCuckoo* builds an extra directed pseudoforest graph structure to present the entire cuckoo graph. Then for each insertion, *SmartCuckoo* checks if the new item will add an edge to an existing cycle, and thus predetermine if the insertion can succeed or not.

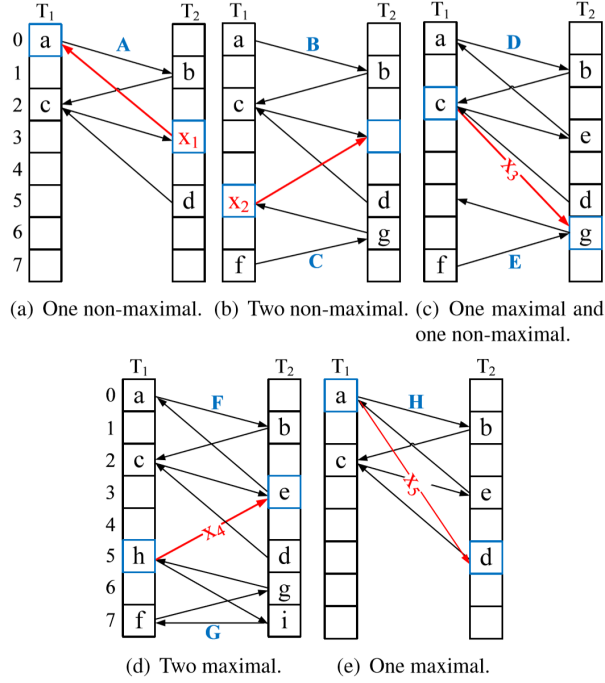
In [SmartCuckoo], a directed graph in which each vertex has an outdegree of exactly one is defined as a **maximal directed pseudoforest**, and a subgraph of the pseudoforest that has the same number of vertices and edges are named **maximal subgraph**. A maximal subgraph contains a cycle, as the outdegree of each vertex is at most one. Fig. 3(a) shows an example of a maximal directed pseudoforest, and there are three maximal subgraphs in the pseudoforest, while Fig. 3(b) shows an example of non-maximal directed pseudoforest which contains two non-maximal subgraphs.




**Fig. 3.** The directed pseudoforest.

In a directed pseudoforest, each vertex corresponds to a bucket, and each edge corresponds to an item and is directed from the its actual store bucket to its alternative bucket. Thus each insertion will add a new edge to the pseudoforest, and add zero, one or two vertices, depending on the two buckets are taken or not. *SmartCuckoo* classifies items insertions into the following three cases based on the number of additional vertices added to the directed pseudoforest, and with help of the directed pseudoforest which tracks the item placements, it can accurately predict the occurrence of endless loops.

- **The Case of  $v+0$ :** When inserting item  $x$ , if the two buckets given by  $h_1(x)$  and  $h_2(x)$  are both occupied, then the insertion will not add new vertex to the pseudoforest. And this case have five possible scenarios, as illustrated by Fig. 4.
  - The two candidate buckets of item  $x_1$ , showns as blue buckets in Fig. 4(a), already exist in the same non-maximal subgraph  $A$ . Either bucket can



**Fig. 4.** Five scenarios for Case  $v + 0$ .



**Fig. 5.** Five scenarios for Case  $v + 1$  and  $v + 2$ .


- be used to place  $x_1$ , as the kick-out operation will finally find an empty bucket in the subgraph.
- The two candidate buckets of item  $x_2$  already exist in two non-maximal subgraph  $B$  and  $C$  respectively, as shown in Fig. 4(b). The insertion will success as both  $B$  and  $C$  have vacant bucket.
  - One candidate bucket of item  $x_3$  is in the non-maximal subgraph  $E$  and the other is in the maximal subgraph  $D$ , as shown in Fig. 4(c).  $x_3$  should be placed in the bucket in the non-maximal subgraph to avoid endless loops.
  - The two candidate buckets of item  $x_4$  already exist in two maximal subgraphs  $F$  and  $G$  as shown in Fig. 4(d). Inserting into either bucket will lead to endless loops.
  - The two candidate buckets of item  $x_4$  already exist in one maximal subgraph  $H$  as shown in Fig. 4(e). Inserting into either bucket will lead to endless loops.
- **The Case of  $v + 1$ :** If one of the candidate bucket of item  $x$  is occupied, i.e., it is in one subgraph, while the other is neither occupied nor a candidate bucket of any other items, inserting  $x$  will add a new vertex to the subgraph, as illustrated by Fig. 5(a).
  - **The Case of  $v + 2$ :** If both the two candidate buckets of item  $x$  have not been claimed by any other items, inserting  $x$  will introduce a new subgraph that formed by the two candidate buckets to the pseudoforest, as illustrated by Fig. 5(b).

[SmartCuckoo] evaluated *SmartCuckoo* and compared it with many other state-of-the-art cuckoo hashing variants—*CHS*, *libcuckoo* and bucketized cuckoo hash tables—in terms of insertion throughput, lookup throughput, and the throughput of mixed queries. The results showed that *SmartCuckoo*’s insertion performance is significantly better than other variants. Specifically, compared to *CHS*, the insertion of *SmartCuckoo* is 0.5 times to 5 times faster, increasing with table load factor, as shown in Fig. 6. The reason is that, when table is heavier loaded, more insertion will be involved in endless loops, and the advantage of skipping those loops becomes more significant.

## 5 Conclusion

Cuckoo hashing introduces a simple hashing scheme that achieves worst-case constant lookup time, which helps storages systems to response to data queries in a real-time manner. However, the kick-out operations of insertion could run into endless loops under unfortunate cases, which makes the insertion performance of cuckoo hashing unpredictable. Many cuckoo hashing variants deal with the endless loops problem by setting an up bound on the number of possible evictions, and perform rehashing if insertion fails. Performing rehashing is expensive, and expected rehashing frequency, although small, may still be unbearable for commercial applications. To address the problem, *CHS* introduces a extra stash





**Fig. 6.** Insertion throughput test.

to place those problematic items that causing the endless loops, and dramatically reduces the probability that a rehashing is required with the cost of a small memory space for stash. *SmartCuckoo*, on the other hand, points out that the loops of kick-out operations during insertion failures are fruitless efforts and waste substantial resources and time, and it is designed to avoid such useless loops completely.

Both *CHS* and *SmartCuckoo* relies on the properties of cuckoo graph to identify endless loops. While *CHS* only tracks a small portion of the graph that involved with new item during insertion, *SmartCuckoo* tracks the entire cuckoo graph and present it using a pseudoforest. The insertion throughput performance of *SmartCuckoo* is much better than *CHS*, however, that performance advantage is traded by a huge memory space cost—extra  $O(n)$  memory is required to keep the pseudoforest. And the difference makes *CHS* and *SmartCuckoo* have different application scenarios: *CHS* is more suitable for system designs where memory space is valuable such as cache designs, while *SmartCuckoo* is better to be applied in the systems where memory is usually sufficient and the performance is a much important issue.

## References

1. Author, F.: Article title. Journal **2**(5), 99–110 (2016)
2. Author, F., Author, S.: Title of a proceedings paper. In: Editor, F., Editor, S. (eds.) CONFERENCE 2016, LNCS, vol. 9999, pp. 1–13. Springer, Heidelberg (2016). <https://doi.org/10.1007/1234567890>
3. Author, F., Author, S., Author, T.: Book title. 2nd edn. Publisher, Location (1999)
4. Author, A.-B.: Contribution title. In: 9th International Proceedings on Proceedings, pp. 1–2. Publisher, Location (2010)
5. LNCS Homepage, <http://www.springer.com/lncs>. Last accessed 4 Oct 2017