

# CIS5516 HW2

Zhijia Chen

## Basic SQL queries:

1. Find the names of all the instructors from Biology department.

**SQL query:**

```
SELECT DISTINCT name FROM instructor WHERE dept_name = 'Biology';
```

**Relational algebra query:**

$$\Pi_{name}(\sigma_{dept\_name="Biology"}(instructor))$$

**Tuple calculus query:**

$$\{t | \exists s \in instructor (t[name] = s[name] \wedge s[dept\_name] = "Biology")\}$$

**Domain calculus query:**

$$\{< name > \mid < id, name, dept, sal > \in instructor \wedge dept = "Biology"\}$$

```
CIS5516=# SELECT DISTINCT name FROM instructor WHERE dept_name = 'Biology';
name
-----
Crick
(1 row)
```

2. Find the names of courses in Computer Science department which have 3 credits.

**SQL query:**

```
SELECT title FROM course WHERE dept_name = 'Comp. Sci.' AND credits = 3;
```

**Relational algebra query:**

$$\Pi_{title}(\sigma_{dept\_name="Comp. Sci." \wedge credits=3}(course))$$

**Tuple calculus query:**

$$\{t | \exists s \in course (t[title] = s[title] \wedge s[dept\_name] = "Comp. Sci." \wedge s[credits] = 3)\}$$

**Domain calculus query:**

$$\{< tit > \mid < cou, tit, dept, cre > \in course \wedge dept = "Comp. Sci." \wedge cre = 3\}$$

```
CIS5516=# SELECT title FROM course WHERE dept_name = 'Comp. Sci.' AND credits = 3;
title
-----
Robotics
Image Processing
Database System Concepts
(3 rows)
```

3. For the student with ID 12345 (or any other value), show all course\_id and title of all courses registered for by the student.

**SQL query:**

```
SELECT course.course_id, title FROM (takes INNER JOIN course ON takes.course_id = course.course_id) WHERE id = '00128';
```

**Relational algebra query:**

$$\Pi_{course\_id, title}(\sigma_{id="00128"}(takes \bowtie course))$$

**Tuple calculus query:**

$$\{t | \exists s \in \text{course}(t[\text{course\_id}] = s[\text{course\_id}] \wedge t[\text{title}] = s[\text{title}] \wedge \exists u \in \text{takes}(s[\text{course\_id}] = u[\text{course\_id}] \wedge u[\text{id}] = "00128"))\}$$

**Domain calculus query:**

$$\{ \langle \text{cou}, \text{tit} \rangle \mid \exists \text{dep}, \text{cre} ( \langle \text{cou}, \text{tit}, \text{dep}, \text{cre} \rangle \in \text{course} \wedge \exists \text{sec}, \text{sem}, \text{yea}, \text{gra} ( \langle "00128", \text{cou}, \text{sec}, \text{sem}, \text{yea}, \text{gra} \rangle \in \text{takes} ) ) \}$$

```
CIS5516=# SELECT course.course_id, title FROM (takes INNER JOIN course ON takes.course_id =
course.course_id) WHERE id = '00128';
course_id |      title
-----+-----
CS-101    | Intro. to Computer Science
CS-347    | Database System Concepts
(2 rows)
```

4. For the student with ID 12345 (or any other value), show the total number of credits taken by that student. Use SQL aggregation on courses taken by the student.

**SQL query:**

```
SELECT sum(credits) FROM (takes INNER JOIN course ON takes.course_id =
course.course_id) WHERE id = '00128';
```

**Relational algebra query:** Not applicable.

**Tuple calculus query:** Not applicable.

**Domain calculus query:** Not applicable.

```
CIS5516=# SELECT sum(credits) FROM (takes INNER JOIN course ON takes.course_id = course.course_id)
WHERE id = '00128';
sum
----
7
(1 row)
```

5. Display the total credits for each student, along with the ID of the student; don't worry about the name of the student. (Don't bother about students who have not registered for any course, they can be omitted).

**SQL query:**

```
SELECT id, sum(credits) FROM (takes INNER JOIN course ON takes.course_id =
course.course_id) GROUP BY id;
```

**Relational algebra query:** Not applicable.

**Tuple calculus query:** Not applicable.

**Domain calculus query:** Not applicable.

```
CIS5516=# SELECT id, sum(credits) FROM (takes INNER JOIN course ON takes.course_id =
course.course_id) GROUP BY id;
```

```
id | sum
-----+-----
76543 | 7
19991 | 3
00128 | 7
98765 | 7
54321 | 8
12345 | 14
45678 | 11
98988 | 8
55739 | 3
76653 | 3
44553 | 4
23121 | 3
(12 rows)
```

6. Find the names of all students who have taken any Comp. Sci. course ever (there should be no duplicate names).

**SQL query:**

```
SELECT DISTINCT name FROM (student INNER JOIN takes ON student.id = takes.id INNER
JOIN course ON takes.course_id = course.course_id) WHERE course.dept_name =
'Comp. Sci.';
```

**Relational algebra query:**

$$\Pi_{name} \left( \sigma_{dept\_name = \text{"Comp. Sci."}} (student \bowtie takes) \right)$$

**Tuple calculus query:**

$$\{t | \exists s \in student (t[name] = s[name] \wedge \exists u \in takes (s[course\_id] = u[course\_id] \wedge u[id] = s[id] \wedge \exists v \in course (v[course\_id] = s[course\_id] \wedge v[dept\_name] = \text{"Comp. Sci."})))\}$$

**Domain calculus query:**

$$\{ \langle nam \rangle \mid \exists id, dep, tot ( \langle id, cou, dep, tot \rangle \in student \wedge \exists sec, sem, yea, gra ( \langle id, cou, sec, sem, yea, gra \rangle \in takes \wedge \exists tit, crd ( \langle cou, tit, \text{"Comp. Sci."}, crd \rangle \in course ) ) ) \}$$

```
CIS5516=# SELECT DISTINCT name FROM (student INNER JOIN takes ON student.id = takes.id INNER JOIN
course ON takes.course_id = course.course_id) WHERE course.dept_name = 'Comp. Sci.';
```

```
name
-----
Bourikas
Brown
Levy
Shankar
Williams
Zhang
(6 rows)
```

7. Display the IDs of all instructors who have never taught a course. (Note: Oracle uses the keyword minus in place of except).

**SQL query:**

```
SELECT id FROM instructor EXCEPT (SELECT id FROM teaches);
```

**Relational algebra query:**

$\Pi_{id}(\sigma(instructor)) - \Pi_{id}(\sigma(teaches))$

**Tuple calculus query:**

$\{t | \exists s \in instructor(t[id] = s[id] \wedge \forall u \in teaches(s[id] \neq u[id]))\}$

**Domain calculus query:**

$\{ \langle id \rangle \mid \exists nam, dep, sal ( \langle id, nam, dep, sal \rangle \in instructor$   
 $\wedge \forall (cou, sec, sem, yea) (\neg \langle id, cou, sec, sem, yea \rangle \in teaches) ) \}$

```
CIS5516=# SELECT id FROM instructor EXCEPT (SELECT id FROM teaches);
```

```
id
-----
33456
58583
76543
(3 rows)
```

## Intermediate SQL queries:

1. Find the maximum and minimum enrollment across all sections, considering only sections that had some enrollment. (Don't worry about those that had no students taking that section.).

**SQL query:**

```
WITH tmp AS (SELECT count(id) AS enrol_num FROM (takes INNER JOIN section ON
takes.course_id = section.course_id AND takes.sec_id = section.sec_id AND takes.year =
section.year) GROUP BY takes.course_id, takes.sec_id, takes.year, takes.semester
HAVING count(id) > 0) SELECT max(enrol_num) AS max_enrol, min(enrol_num) AS
min_enrol FROM tmp;
```

**Relational algebra query:** Not applicable.

**Tuple calculus query:** Not applicable.

**Domain calculus query:** Not applicable.

```
CIS5516=# WITH tmp AS (SELECT count(id) AS enrol_num FROM (takes INNER JOIN section ON
takes.course_id = section.course_id AND takes.sec_id = section.sec_id AND takes.year = section.year)
GROUP BY takes.course_id, takes.sec_id, takes.year, takes.semester HAVING count(id) > 0) SELECT
max(enrol_num) AS max_enrol, min(enrol_num) AS min_enrol FROM tmp;
```

```
max_enrol | min_enrol
-----+-----
6 | 1
(1 row)
```

2. Find all sections that had the maximum enrollment (along with the enrollment), using a subquery.

**SQL query:**

```
WITH tmp AS (SELECT count(id) AS enrol_num, takes.year, takes.semester,
takes.course_id, takes.sec_id FROM (takes INNER JOIN section ON takes.course_id =
section.course_id AND takes.sec_id = section.sec_id AND takes.year = section.year)
GROUP BY takes.course_id, takes.sec_id, takes.year, takes.semester HAVING count(id) >
0) SELECT year, semester, course_id, sec_id, enrol_num FROM tmp WHERE enrol_num =
(SELECT max(enrol_num) FROM tmp);
```

**Relational algebra query:** Not applicable.

**Tuple calculus query:** Not applicable.

**Domain calculus query:** Not applicable.

```
CIS5516=# WITH tmp AS (SELECT count(id) AS enrol_num, takes.year, takes.semester, takes.course_id,
takes.sec_id FROM (takes INNER JOIN section ON takes.course_id = section.course_id AND takes.sec_id =
section.sec_id AND takes.year = section.year) GROUP BY takes.course_id, takes.sec_id, takes.year,
takes.semester HAVING count(id) > 0) SELECT year, semester, course_id, sec_id, enrol_num FROM tmp
WHERE enrol_num = (SELECT max(enrol_num) FROM tmp);
year | semester | course_id | sec_id | enrol_num
-----+-----+-----+-----+-----
2009 | Fall    | CS-101   | 1      | 6
(1 row)
```

3. Find all courses whose identifier starts with the string "CS-1".

**SQL query:**

```
SELECT title FROM course WHERE course_id LIKE 'CS-1%';
```

**Relational algebra query:**

$$\Pi_{title} \left( \sigma_{course\_id = "CS-1\%"}(course) \right)$$

**Tuple calculus query:**

$$\{t | \exists s \in course (t[title] = s[title] \wedge s[course\_id] = "CS-1\%")\}$$

**Domain calculus query:**

$$\{ \langle tit \rangle \mid \langle cou, tit, dep, sal \rangle \in course \wedge cou = "CS-1\%" \}$$

```
CIS5516=# select title from course where course_id like 'CS-1%';
title
-----
Intro. to Computer Science
Game Design
(2 rows)
```

## Advanced SQL queries:

1. Create a view faculty showing only the ID, name, and department of instructors.

**SQL query:**

```
CREATE VIEW faculty AS SELECT id, name, dept_name FROM instructor;
```

**Relational algebra query:**

$$\rho_{faculty} \left( \Pi_{id, name, dept\_name} \left( \sigma(instructor) \right) \right)$$

**Tuple calculus query:** Not applicable.

**Domain calculus query:** Not applicable.

```
CIS5516=# CREATE VIEW faculty AS SELECT id, name, dept_name FROM instructor;
CREATE VIEW
CIS5516=# select * from faculty limit 1;
id | name | dept_name
-----+-----+-----
10101 | Srinivasan | Comp. Sci.
(1 row)
```

2. Create a view CSInstructors, showing all information about instructors from the Comp. Sci. department.

**SQL query:**

```
CREATE VIEW CSInstructor AS SELECT * FROM instructor WHERE dept_name = 'Comp. Sci.';
```

**Relational algebra query:**

$\rho_{CSInstructor}(\sigma_{dept\_name="Comp. Sci."}(instructor))$

**Tuple calculus query:** Not applicable.

**Domain calculus query:** Not applicable.

```
CIS5516=# CREATE VIEW CSInstructor AS SELECT * FROM instructor WHERE dept_name = 'Comp. Sci.';
CREATE VIEW
CIS5516=# SELECT * FROM CSInstructor LIMIT 1;
 id | name | dept_name | salary
-----+-----+-----+-----
10101 | Srinivasan | Comp. Sci. | 65000.00
(1 row)
```

3. Find all rooms that have been assigned to more than one section at the same time. Display the rooms along with the assigned sections. (Hint: use WITH or views.)

**SQL query:**

```
SELECT course_id, sec_id, room_number, building FROM section WHERE (building, room_number, time_slot_id) IN (SELECT building, room_number, time_slot_id FROM section GROUP BY building, room_number, time_slot_id HAVING count(course_id) > 1);
```

**Relational algebra query:** Not applicable.

**Tuple calculus query:** Not applicable.

**Domain calculus query:** Not applicable.

```
CIS5516=# select course_id, sec_id, room_number, building from section where (building, room_number, time_slot_id) in (select building, room_number, time_slot_id from section group by building, room_number, time_slot_id having count(course_id) > 1);
 course_id | sec_id | room_number | building
-----+-----+-----+-----
CS-190    | 2      | 3128        | Taylor
CS-319    | 2      | 3128        | Taylor
CS-347    | 1      | 3128        | Taylor
EE-181    | 1      | 3128        | Taylor
(4 rows)
```